

xmm : eXtended Matrix Multiplication

1. Formulation

xmm generalizes the matrix multiplication operation to a customizable form:

$$Z = \Omega_{\{\sigma, \mu\}}(A_1, \dots, A_{N_a}, B_1, \dots, B_{N_b})$$
$$Z(i, j) \stackrel{\text{def}}{=} \sigma_k \left\{ \mu(A_1(i, k), \dots, A_{N_a}(i, k), B_1(k, j), B_{N_b}(k, j)) \right\}$$

where operation Ω is parameterized by two functions, μ the **combinator**, and σ the **reductor**.

2. Relations to neural networks

2.1. Matrix Multiplication

matmul is a special case of xmm where

- $\sigma = \sum$,
- $\mu(x, y) = xy$,
- $N_a = N_b = 1$.

2.2. Kolmogorov Arnold Network

KAN-like operations is a special case of xmm where

- $\sigma = \sum$
- $\mu(r, c_1, \dots, c_n)$ is a customizable term.
 - Take polynomial KAN with `gridsize = 4` as an example.

$$\mu(r, c_1, \dots, c_n) = c_1 + c_2 r + c_3 r^2 + c_4 r^3$$

- $N_a = 1, N_b = \text{gridsize}$.

3. General Interface Description

xmm should be able to:

- Provide interface to customize μ . This *could* be done by:
 - parse an expression into μ .
 - other possible ways...

- Further, provide interface to customize σ by parsing a binary function $\mathbb{R}^2 \rightarrow \mathbb{R}$ into σ , resembling passing a functor to the `std::reduce` function in C++ STL.
- Just-in-time compile the operator for runtime efficiency.
- **Support pytorch autograd and computations on CUDA**, to provide gradient to **all** input tensors.
- **Provide optimized CUDA operator implementations.**

3.1. Side Notes

- I guess customizing μ is of higher priority than σ since I have not yet encountered scenario where σ should be customized to be other than \sum .
- B (column operands) should be transposed to comply with the PyTorch conventions, i.e. the input shapes are (M, K) for row operands and (N, K) for column operands.

4. Interface

4.1. Class SumOperator

- An operator with summation as its **reductor**, i.e. $\sigma = \sum$.

4.1.1. SumOperator.__init__(self, nr: int, nc: int, expr: str)

- Create **and parse** the operator, including:
 - Parse expr;
 - Compute (symbolic) derivatives;
 - Create kernels sources (codegen, forward and backward);
 - Save kernel identifier.

4.1.2. SumOperator.compile(self, build_dir: Optional[str] = None, identifier:

Optional[str] = None):

- Compile the operator with `torch.utils.cpp_extensions`.
- Use an identifier (a str which has pattern `^[A-Za-z0-9_]+$`) to identify different operators.
First If identifier is not specified, a uuid is generated for each new compilations.

- *Simply use the default parameters to recompile before every run.*

4.1.3. `SumOperator.forward(self, *args)`

- Execute the operation defined by the operator.
- `*args` should be row operands followed by column operands, in correct order.

4.1.4. `SumOperator.backward(self, grad_out: torch.Tensor, *args)`

- Compute derivatives w.r.t. all row operands and column operands.
- `grad_out` should be of shape (M, N) .
- `*args` should be row operands followed by column operands, in correct order.

4.1.5. Operator syntax

Take `nr = 1, nc = 3` as an example.

- All identifier you can use is `r1` for the row operand, and `c1, c2, c3` for the three column operands respectively, where `r1 = R1[i, k]`, `c1 = C1[j, k]`.
- Allow a fixed set of simple function, like `exp, sin, tanh, atan2`, etc.
- Constants should only contain its value, **not** followed by `f` or any other suffixes.

A legal expression is like:

```
"c1 * (c2 * c2 * (r1 + c3) * (r1 + c3) - 1.0) * exp(-0.5 * (r1 * r1 * c2 * c2))"
```

5. Modules

5.1. Module `xmm.preprocess`

5.1.1. `preprocess.expr2ast(expr: str)`

5.1.2. `preprocess.ast2sympy(ast_tree)`

5.1.3. `preprocess.sympy2ast(sympy_expr)`

5.1.4. `preprocess.ast2CUDAexpr(ast_rootnode)`

5.2. Module `xmm.codegen`

Generates the CUDA kernels to be compiled.

5.2.1. `codegen.generate_operator_source(nrow: int, ncol: int, expression) ->`

`Tuple[str, str]`

- Returns both C++ wrapper and CUDA source code.

6. Example

We take the following expression as an example.

$$z(i, j) = \sum_k c_1(k, j) \cdot x(i, k) + c_2(k, j) \cdot x^2(i, k) + c_3(k, j) \cdot x^3(i, k) + b(k, j)$$

where $\sigma_k\{\cdot\} = \sum_k$. APIs *could* be like:

```
from xmm.SumOperator import SumOperator
import torch
import torch.nn as nn

op = SumOperator(1, 3, "0.867325070 * c3 * (c1 * c1 * (r1 + c2) * (r1 + c2) - 1)
* exp(-0.5 * c1 * c1 * (r1 + c2) * (r1 + c2))")
op.compile(identifier="mexhat_op")

class XmmFn(torch.autograd.Function):
    @staticmethod
    def forward(ctx, x, scale, bias, weight):
        ctx.save_for_backward(x, scale, bias, weight)
        return op.forward(x, scale, bias, weight)

    @staticmethod
    def backward(ctx, grad_output):
        x, scale, bias, weight = ctx.saved_tensors
        grad_x, grad_scale, grad_bias, grad_weight = op.backward(grad_output, x,
scale, bias, weight)
        return grad_x, grad_scale, grad_bias, grad_weight

class XmmLayer(nn.Module):
    def __init__(self, in_features, out_features):
        super().__init__()
        self.in_features = in_features
        self.out_features = out_features
        self.scale = nn.Parameter(torch.ones(out_features, in_features))
```

```

self.bias = nn.Parameter(torch.zeros(out_features, in_features))
self.weight = nn.Parameter(torch.empty(out_features, in_features))
nn.init.kaiming_uniform_(self.weight, a=(5 ** .5))
self.base_activation = nn.SiLU()
self.bn = nn.BatchNorm1d(out_features)

def forward(self, x):
    x = XmmFn.apply(x, self.scale, self.bias, self.weight)
    return self.bn(x)

```